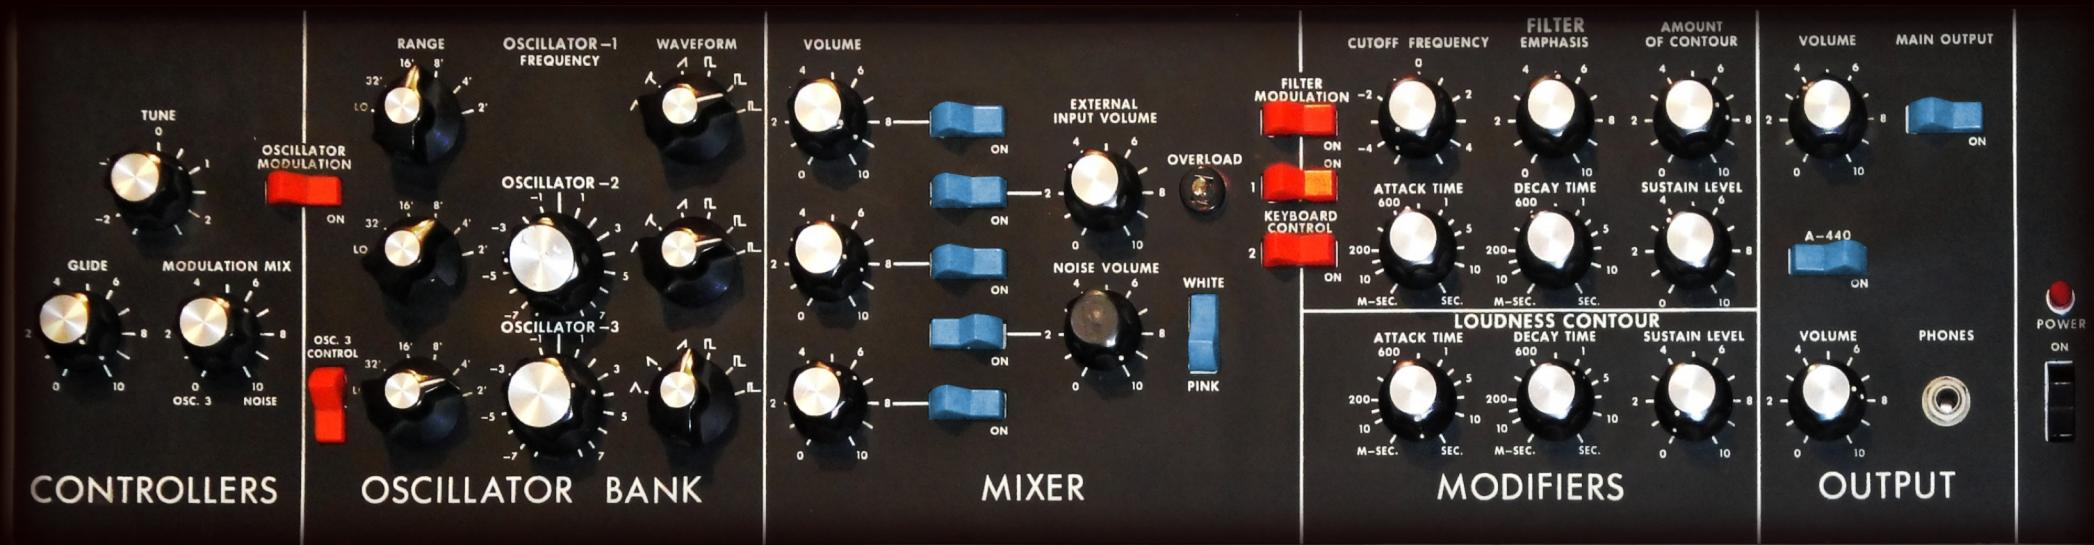


Lecture 4: Subtractive Synthesis



ADDITIVE vs. SUBTRACTIVE SYNTHESIS

Additive Synthesis	Subtractive Synthesis
Powerful	Creation of precise timbres more difficult
Many input parameters	Fewer input parameters
User modulation difficult	User modulation easier
Potentially CPU intensive	CPU efficient
Creation of rich sounds requires much work	Rich sounds are created easily

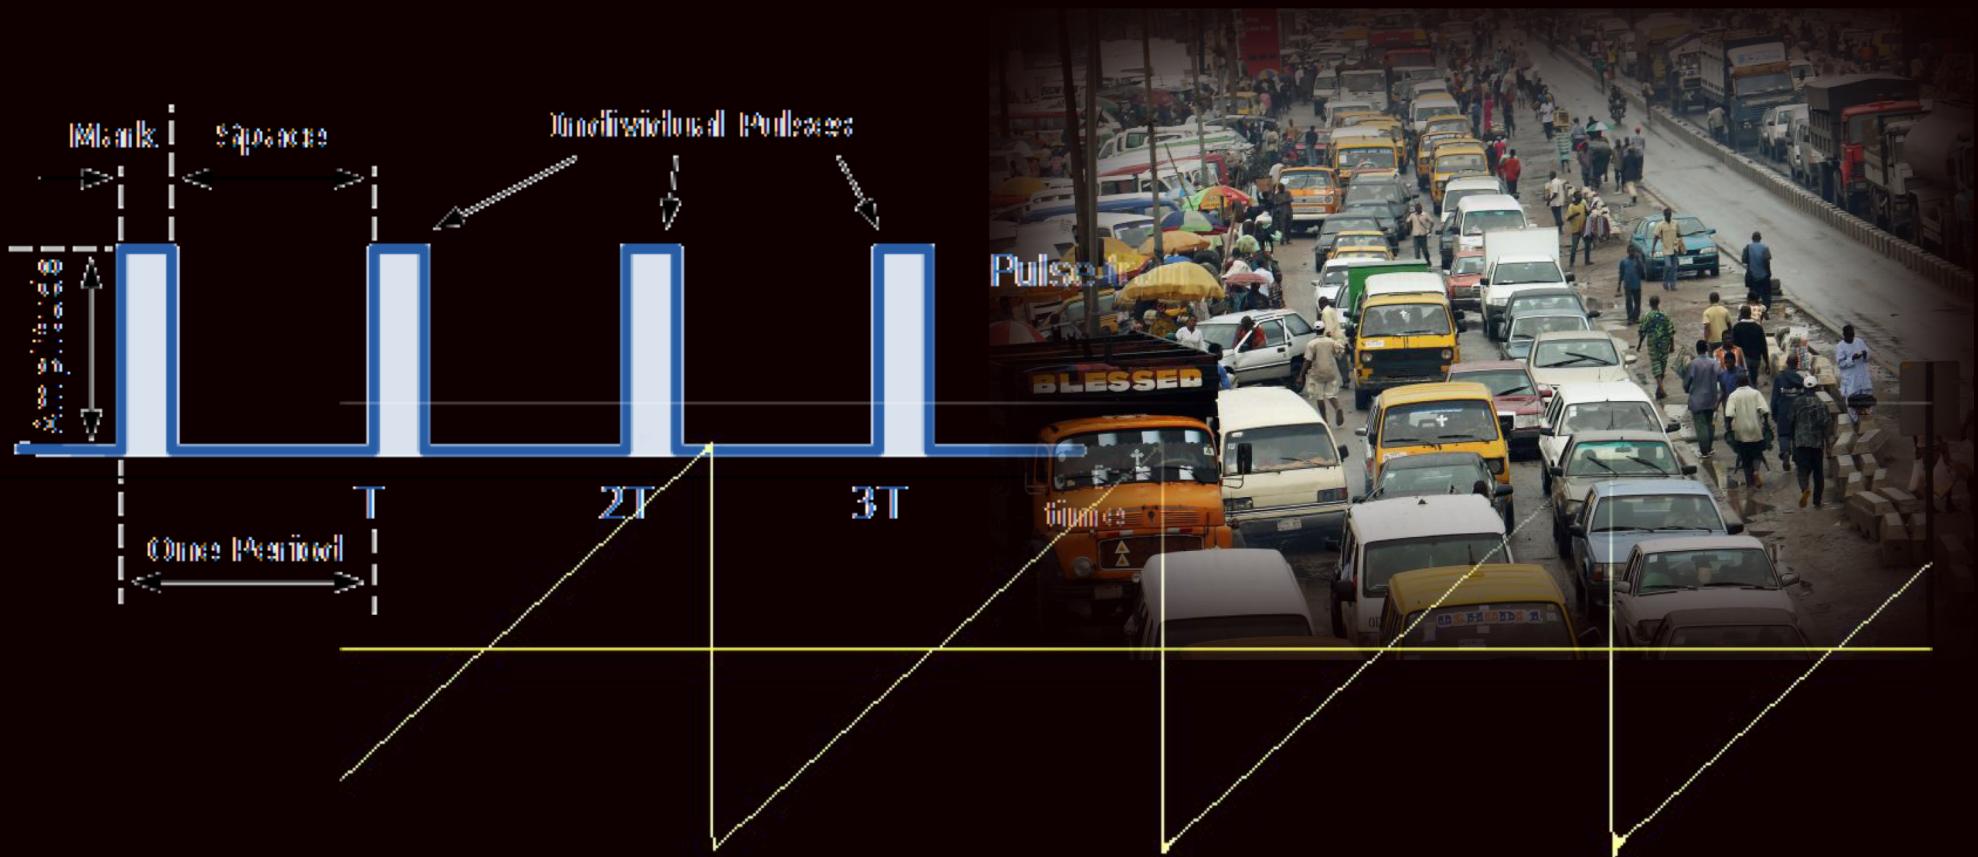
SUBTRACTIVE SYNTHESIS DEFINITION

- Additive synthesis:
 - Based on the notion of building up complex sounds from simple 'atomic' sound elements.
- Subtractive synthesis:
 - Takes the opposite approach by beginning with a complex sound and reducing and refining it through filtering.



RICH SOUND SOURCES

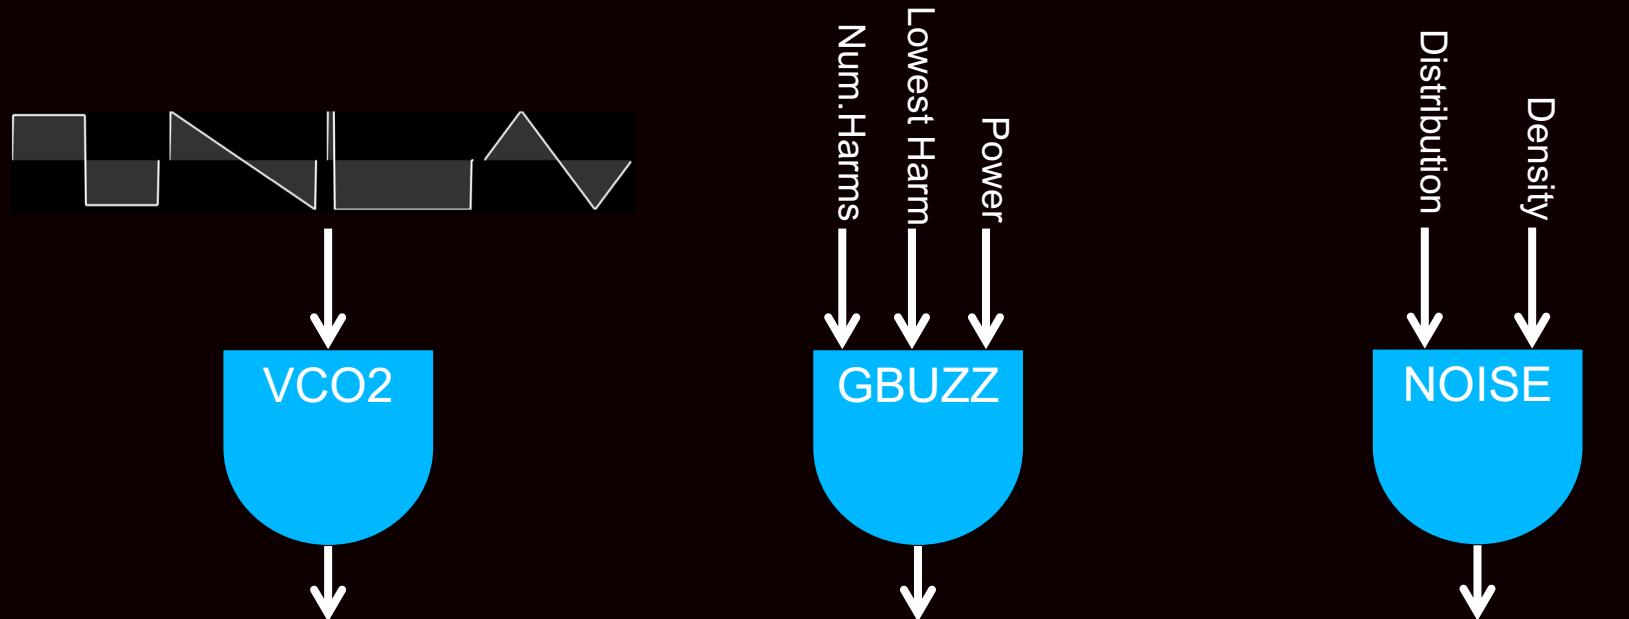
- Our rich sound source for subtractive synthesis could be a sawtooth waveform, a pulse waveform, a square waveform, generated noise, a sound sample of the sea, someone speaking, the traffic...



- A sine tone or a tuning fork are not rich sound sources.

RICH SOUND SOURCES IN CSOUND

- **vco2, squinewave**
 - Useful for generating band-limited basic waveshapes: sawtooth, square, triangle, pulse etc. (example 1 & 2)
- **gbuzz**
 - Useful for creating dynamic spectra based on stacks of harmonically related cosine waves. (example 3)
- **noise, pinkish, pinker, dust2, gausstrig...**
 - Useful for creating noisy source sounds. (example 4)



VCO2

- Opcode format

```
aSig vco2 kAmp, kCPS, iTYPE, kW
```

```
; sawtooth waveform  
aSig vco2 0.2, 220
```

- **iType = 0** produces a sawtooth waveform but this is also the default, so can be omitted.

```
; square waveform  
aSig vco2 0.2, 220, 2, 0.5
```

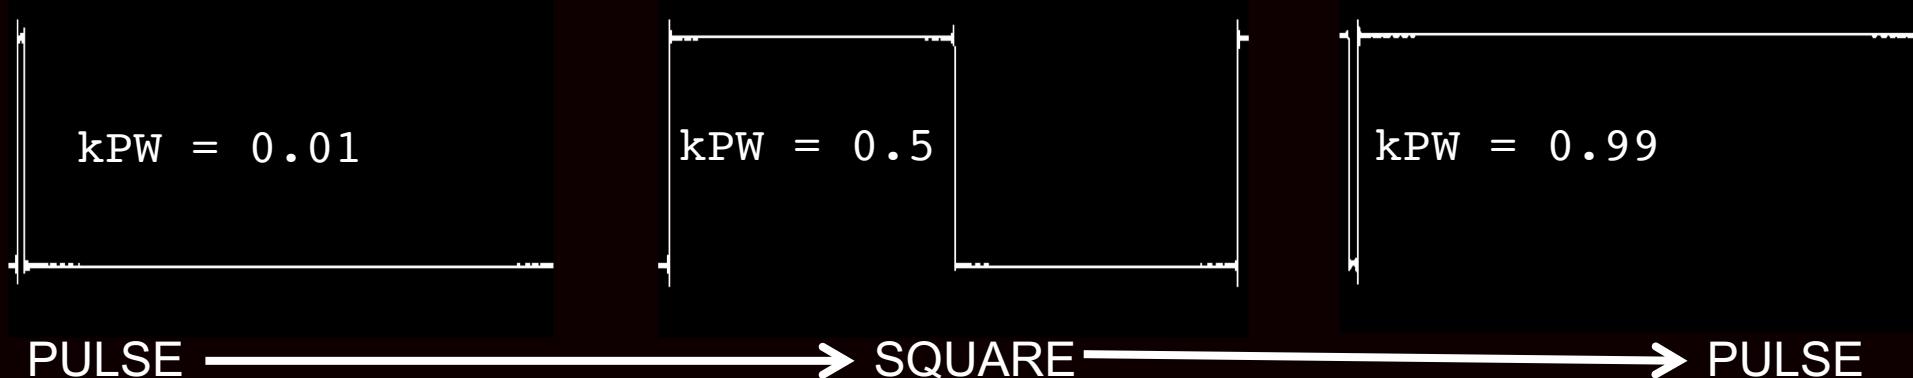
```
; triangle waveform  
aSig vco2 0.2, 220, 4, 0.5
```

```
; pulse waveform  
aSig vco2 0.2, 220, 2, 0.05
```

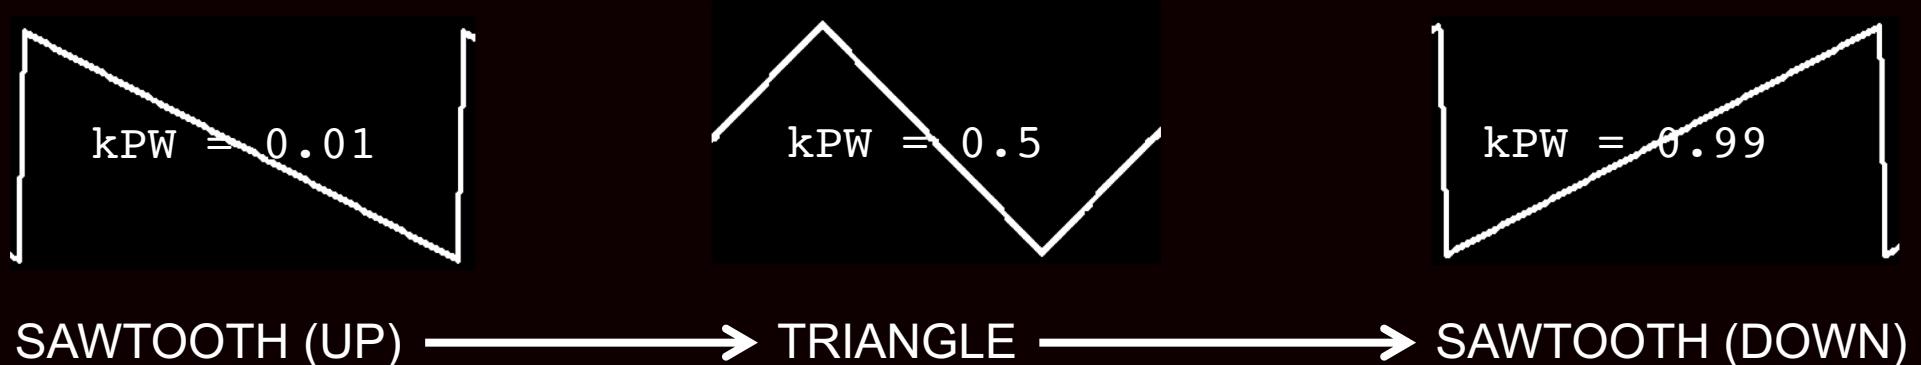
- See example 1.

VCO2 – PULSE-WIDTH MODULATION

```
; type = 2  
kPW    line    0.01, p3, 0.99  
aSig   vco2    0.2, 220, 2, kPW
```



```
; type = 4  
kPW    line    0.01, p3, 0.99  
aSig   vco2    0.2, 220, 4, kPW
```

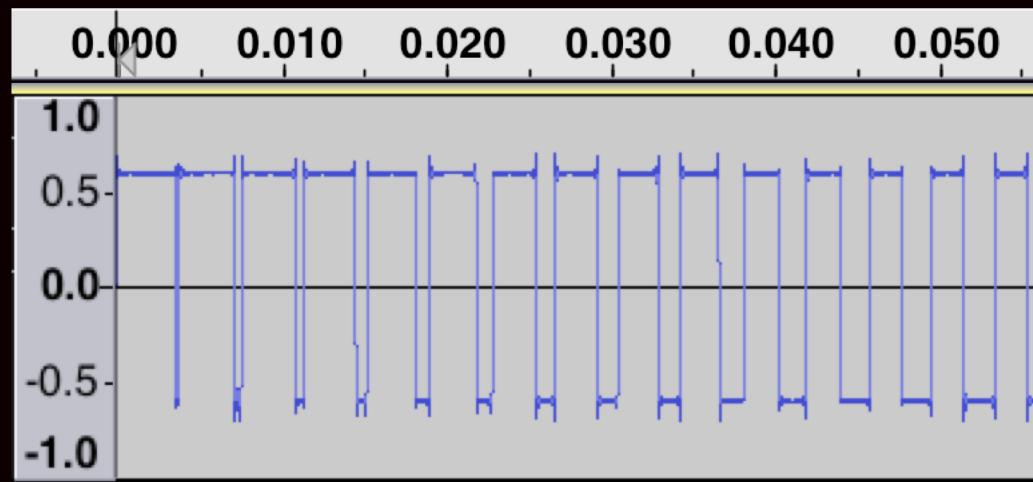


VCO2 – PWM ENVELOPE

- Envelope modulation of pulse-width using an envelope can be used to timbrally articulate a sound.

```
kpw  expseg  0.97, 0.2, 0.5, p3-0.2, 0.5  
asig vco2    0.6, 220, 4, kpw
```

- In the past this has been used to imitate the articulation (attack) of brass instruments but this approach and sound is a little dated nowadays. (example 2)



squinewave

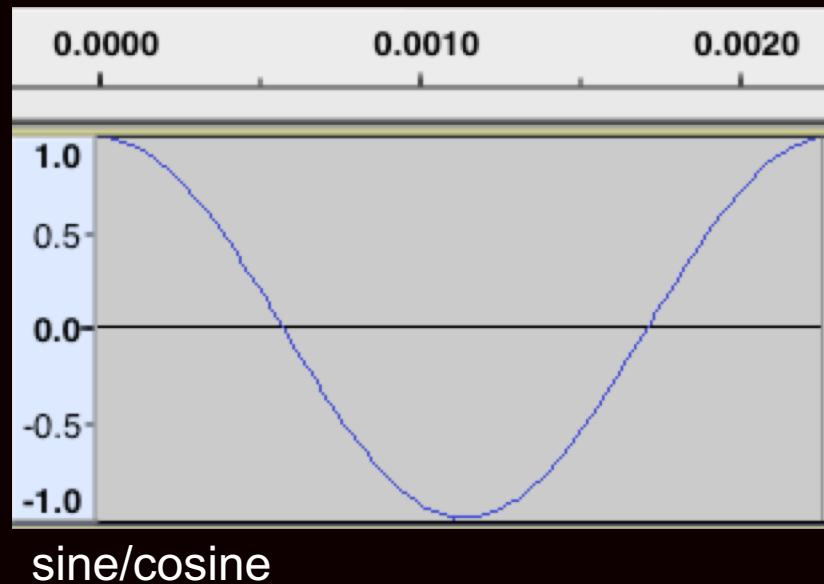
- **squinewave** creates a waveform by joining three segments: the first half of cosine, a flat portion and the second half of a cosine wave.
- By varying the proportions of these three segments we can morph between sine, sawtooth, square and pulse waveforms

```
aout [, asyncout] squinewave acps, aClip, aSkew [,  
asyncin] [, iMinSweep] [, iphase]
```

- The controls we will consider here will be the **aclip** argument, which controls the proportion of the flat portion, and **aSkew** which weights the proportions of the first half and second half of the cosine wave.
- Both these controls need to be a-rate.
- We will not consider squinewave's **syncout** and **syncin** options
- **iMinSweep** specifies a restriction on sudden changes in amplitude in the waveform – this can therefore be used in bandlimiting and as a soft tone control.

squinewave

```
;  
asig, as  squinewave  a(icps), a(0), a(0), a(0)  clip  skew  syncin
```

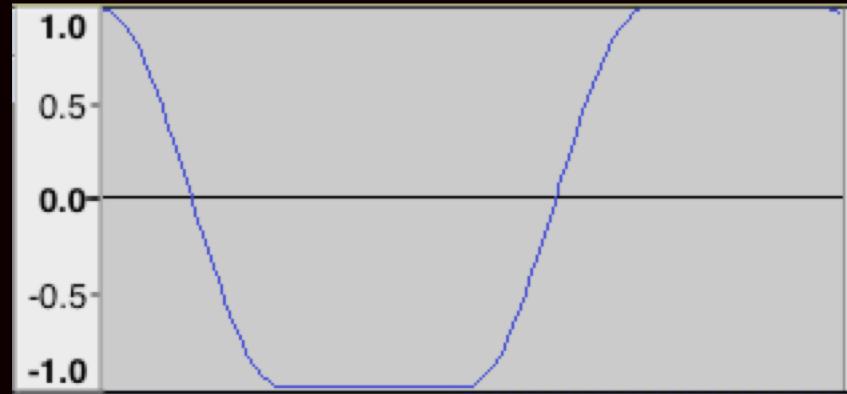


- A simple sinusoidal shape (cosine).
- N.B. **asyncin** is not an optional argument, as suggested in the documentation

squinewave

- Increasing clip above zero increases the proportion of the flat portion.

```
;  
asig, as squinewave clip skew syncin  
a(icps), a(0.5), a(0), a(0)
```



```
;  
asig,as squinewave clip skew syncin minsweep  
a(icps), a(0.5), a(0), a(0), 4
```

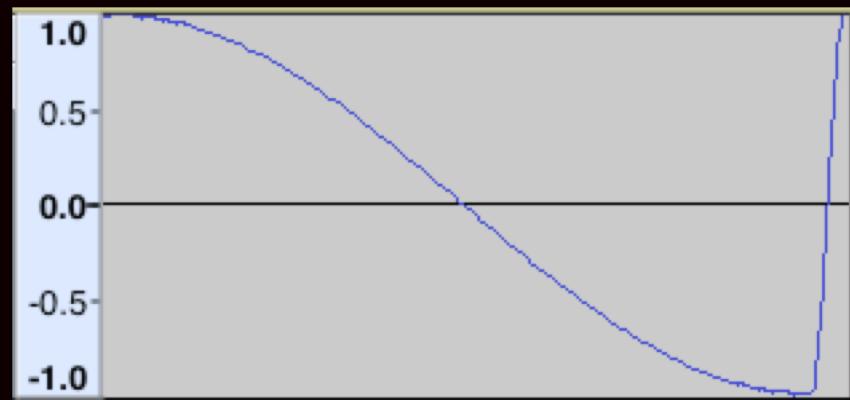


‘Square’ wave

squinewave

- `skew` values closer to -1 and +1 allow access to sawtooth and pulse waveforms.

```
;  
asig, as squinewave clip skew syncin  
a(icps), a(0), a(-1), a(0)
```



‘Sawtooth’ waveform

```
;  
asig,as squinewave clip skew syncin minsweep  
a(icps), a(1), a(-1), a(0), 4
```



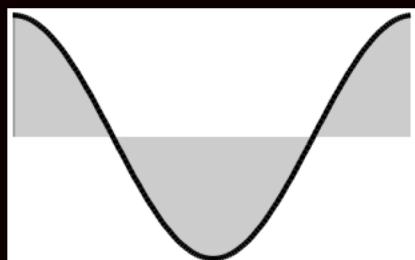
‘Pulse’ waveform

GBUZZ

- **gbuzz** creates dynamic spectra based on stacks of harmonically related cosine waves.

```
gicos ftgen    0,0,131072,11,1 ; a cosine wave  
  
; snip  
  
knh    =          30           ; number of harmonics  
klh    =          1            ; lowest harmonic  
kmul  line     0, p3*0.5, 0.9, p3*0.5, 0  
                  ; power multiplier  
kmul  gbuzz    kamp, icps, knh, klh, kmul, gicos
```

- **gbuzz** requires us to supply it with a cosine wave. We can do this with GEN 11.



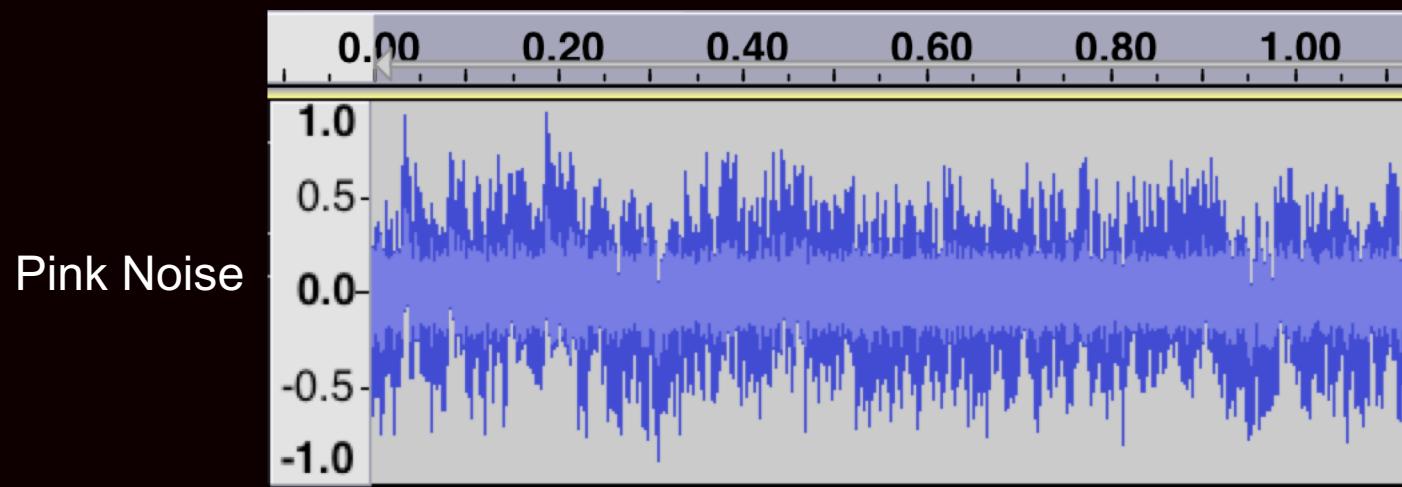
Cosine Wave

GBUZZ

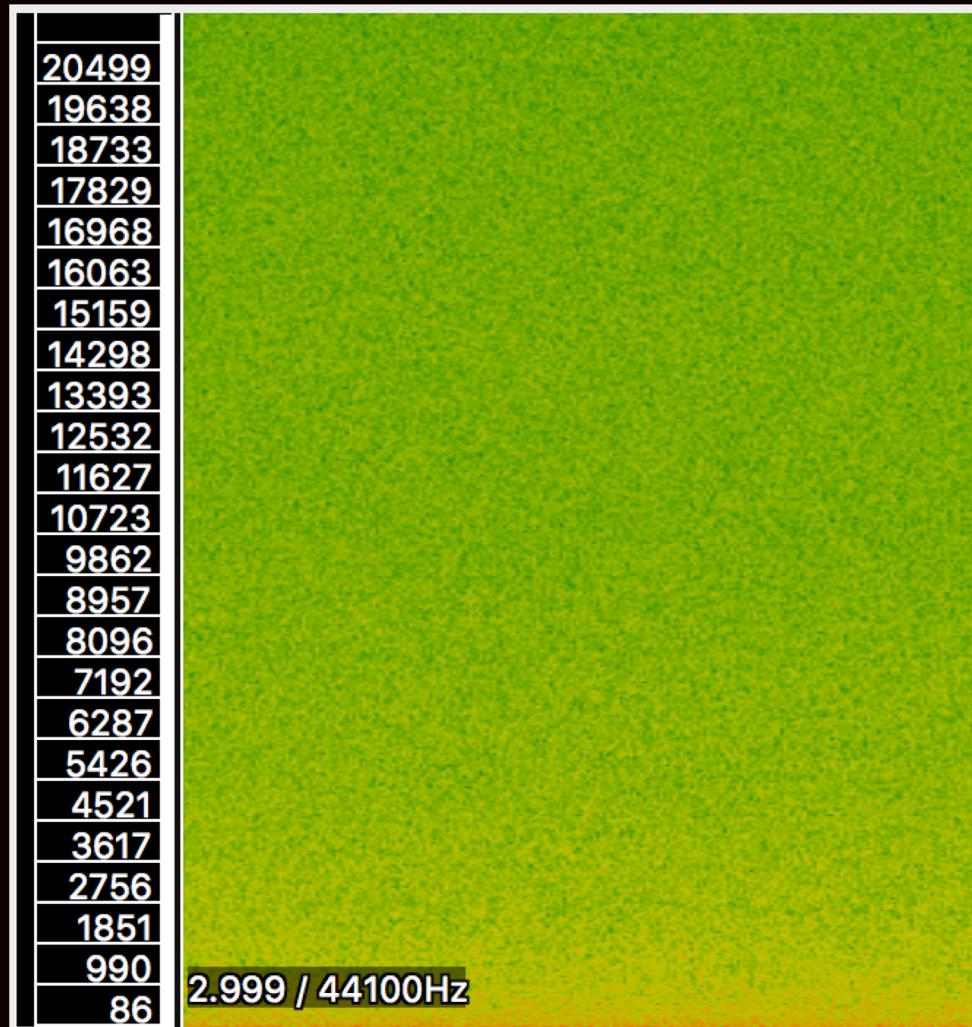
- **knh** (number of harmonics) defines the extent of the harmonic spectrum in harmonic partials above the lowest partial.
- **k1h** (lowest harmonic) is the lowest harmonic partial to include in the stack. To start at the fundamental this should be 1.
- Although **knh** and **k1h** can be changed at k-rate, they actually only change in integer steps, therefore clicks can occur when partials enter or leave.
- **kmul** controls the amplitude weighting of the partials:
 - values zero to 1 favour lower partials.
 - 1 results in all partials being equal in strength.
 - values greater than 1 favour the upper partials.
- Modulating **kmul** can be used to produce dynamic spectra.
- See example 3.

VARIETIES OF 'NOISE'

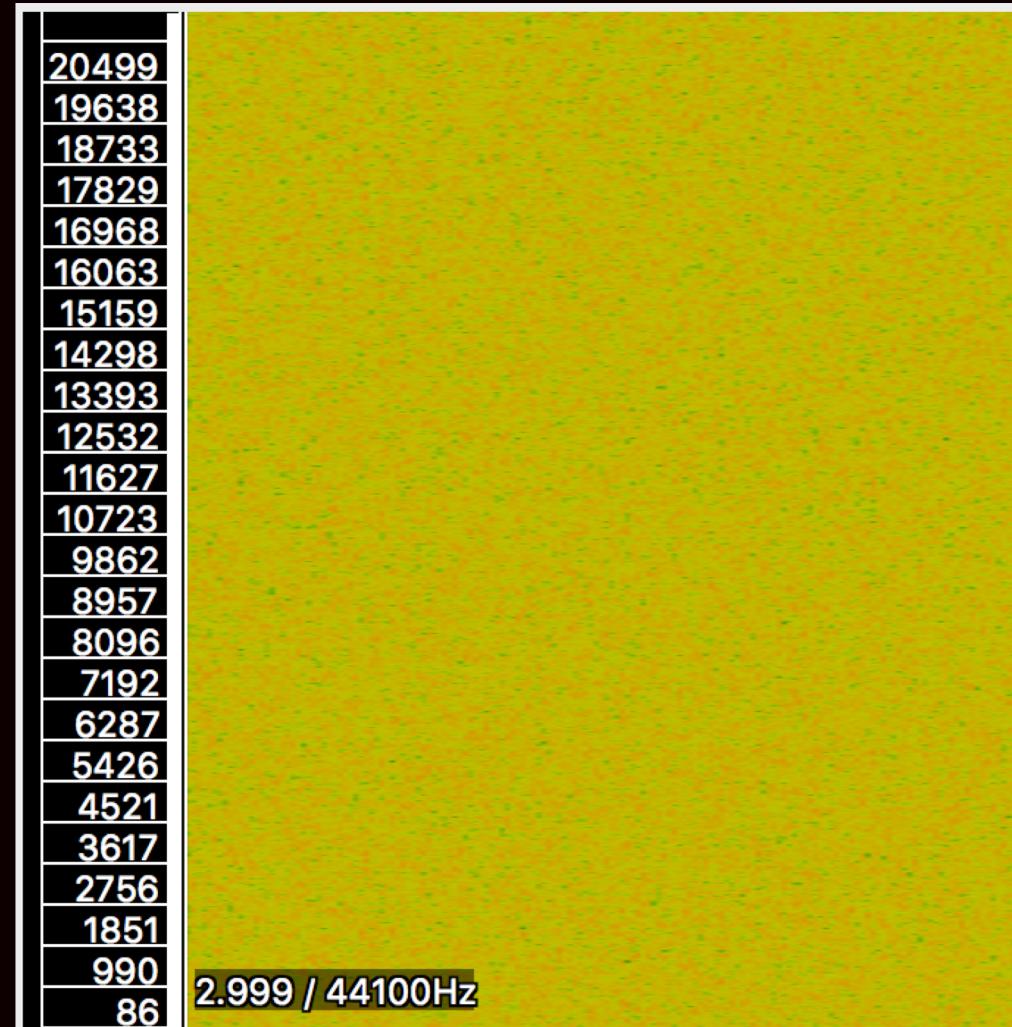
- 'Noise' can be dull or unpleasant on its own but offers possibilities for refinement through subtractive synthesis.
- 'White' noise presents random noise with equal energy at all frequencies. However we perceive frequency with a logarithmic scale so this is not considered 'musical' and can sound excessively bright and shrill.
- 'Pink' noise presents random noise with equal energy in all octaves. This can be considered more 'musical' therefore is normally more useful as a sound source in subtractive synthesis.
- 'Dust' can be considered noise, wherein the noise particles are interspersed with gaps of silence. Differing densities of dust can be given poetic descriptions such as 'wind', 'sand', 'gravel', 'grit' etc.



VARIETIES OF 'NOISE'



Pink Noise



White Noise

NOISE OPCODES

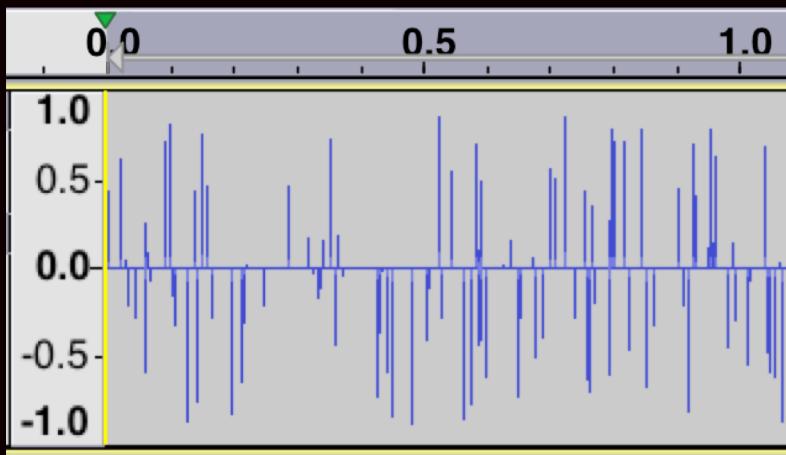
```
kBeta      =      0  
aWhiteNoise noise  kAmp, kBeta
```

- **kBeta** controls an internal lowpass filter but for the generation of white noise this should be zero.

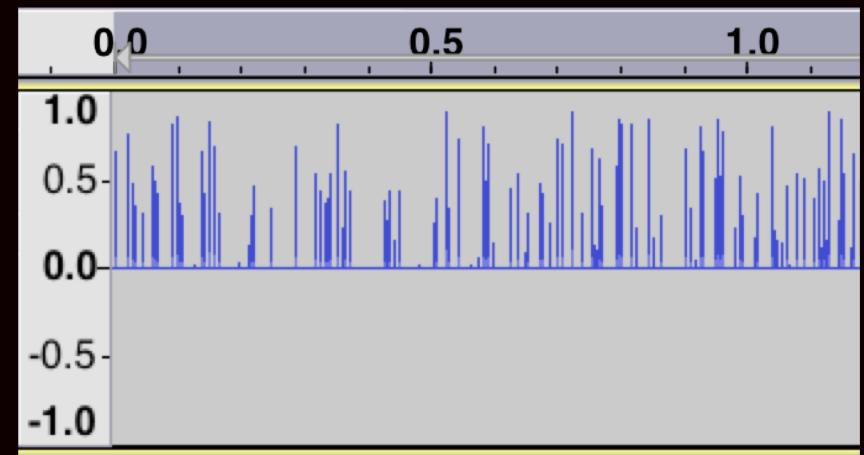
```
aPinkNoise pinkish kAmp
```

```
aDust   dust2   kAmp, kDensity
```

- **dust2** produces bipolar noise, dust produces unipolar noise.



Bipolar Noise (dust2)

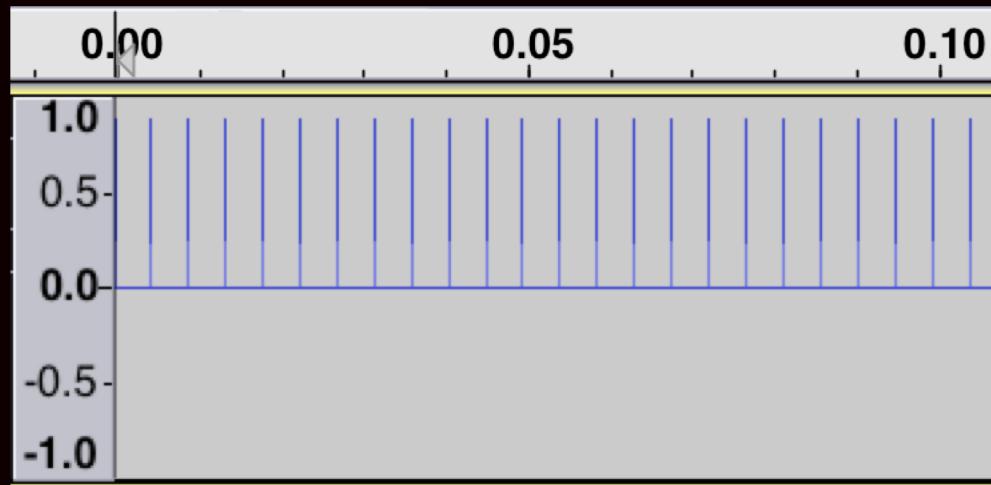


Unipolar Noise (dust)

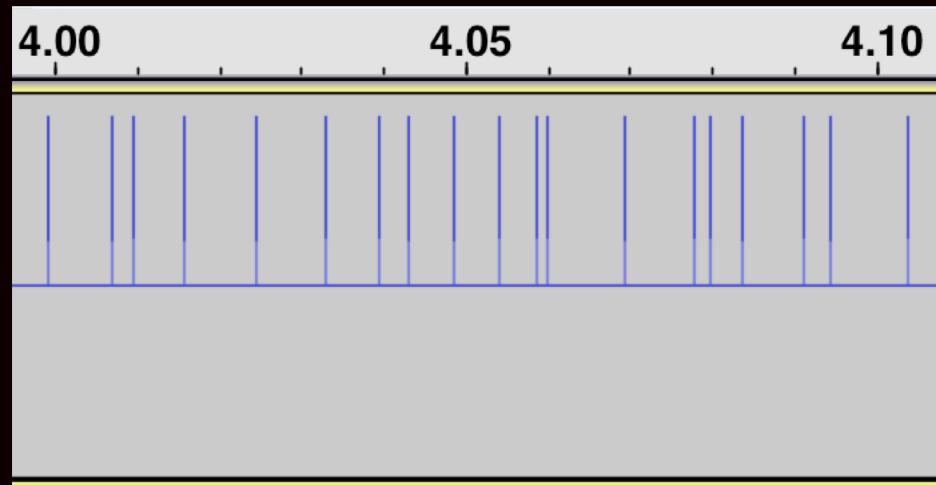
NOISE OPCODES - gausstrig

```
aSig  gausstrig  kAmp,  kCPS,  kDev
```

- With **gausstrig** we can move between periodic noise (a thin buzzing sound) and aperiodic noise by varying its **kDev** parameter.
- Its output is unipolar.
- Density of the noise or ‘pulses’ is controlled by **kCPS**.
- Pulse amplitudes could also be randomly scattered by controlling **kAmp** with a random function.



Periodic (kDev=0, kCPS=220)



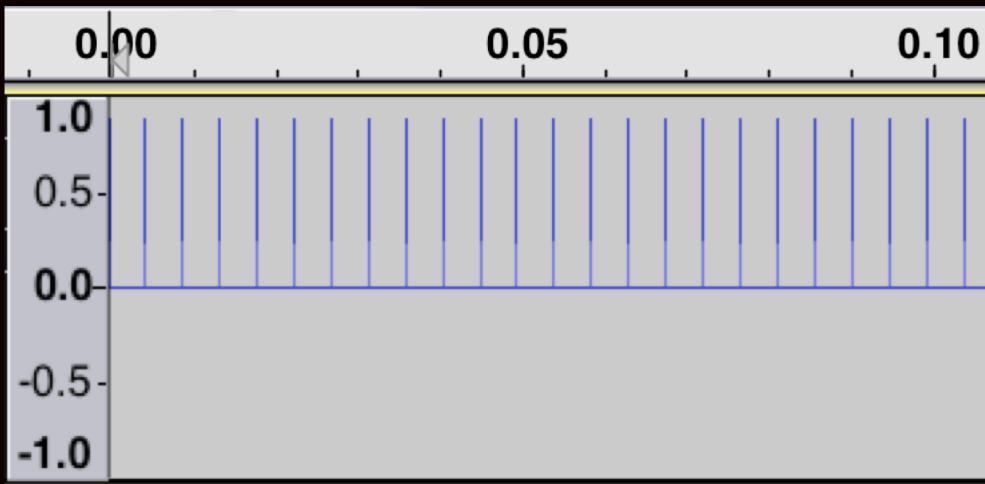
Aperiodic (kDev=1, kCPS=220)

- Example 4

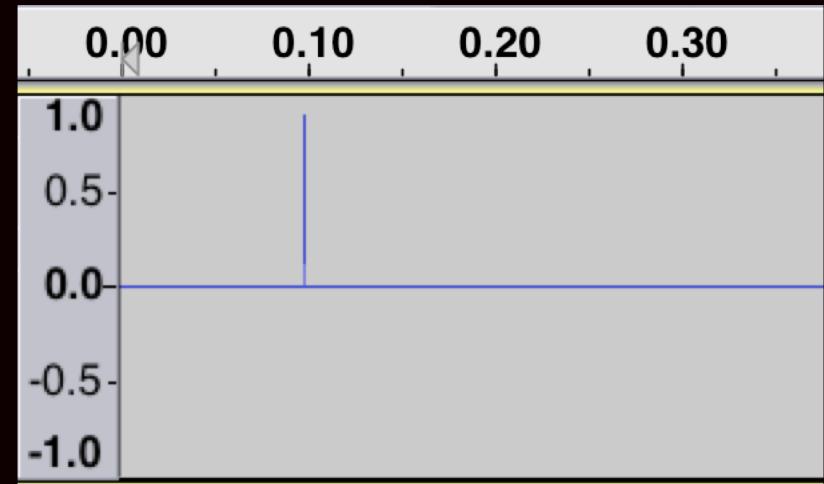
NOISE OPCODES - mpulse

```
aSig mpulse kAmp, kGap
```

- **mpulse** creates a train of impulses similar to **gausstrig** with no deviation applied.
- Instead of specifying a frequency we specify the time gap (**kGap**) between impulses. Therefore we can derive the resulting frequency as $1/\text{time_gap}$.
- We have a special setting: if we set time gap to zero then **mpulse** will then just produce a single impulse and then nothing more. This can be useful when using subtractive synthesis for creating percussive sounds.



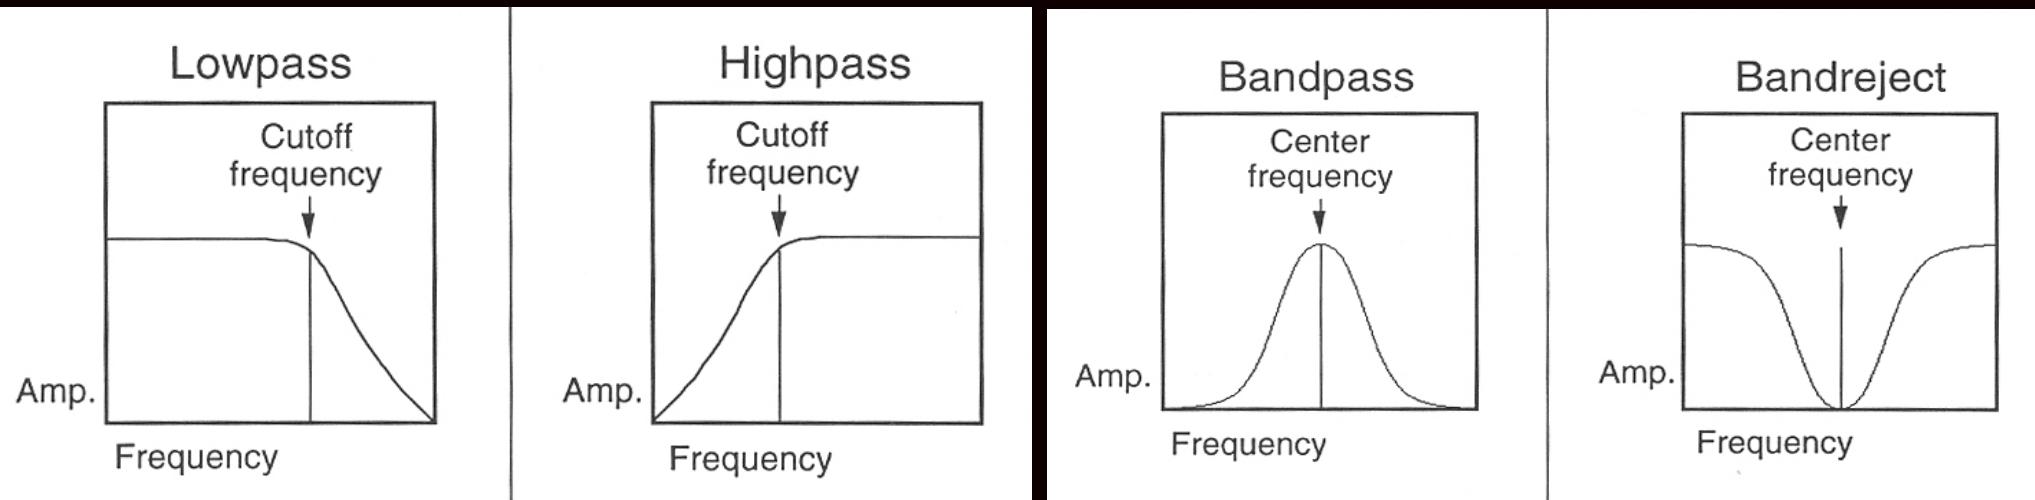
A train of impulses



A single impulse

BASIC FILTERS

- The four common filter types are:
 - Lowpass
 - Highpass
 - Bandpass
 - Bandreject (notch)
- **Lowpass** (or **highcut**) allows frequencies below its cutoff frequency to pass unattenuated.
- **Highpass** (or **lowcut**) allows frequencies above its cutoff frequency to pass unattenuated.
- **Bandpass** allows frequencies at and near its cutoff frequency to pass.
- **Bandreject** (or **notch**) stops frequencies at and around its cutoff frequency.



CUTOFF SLOPES

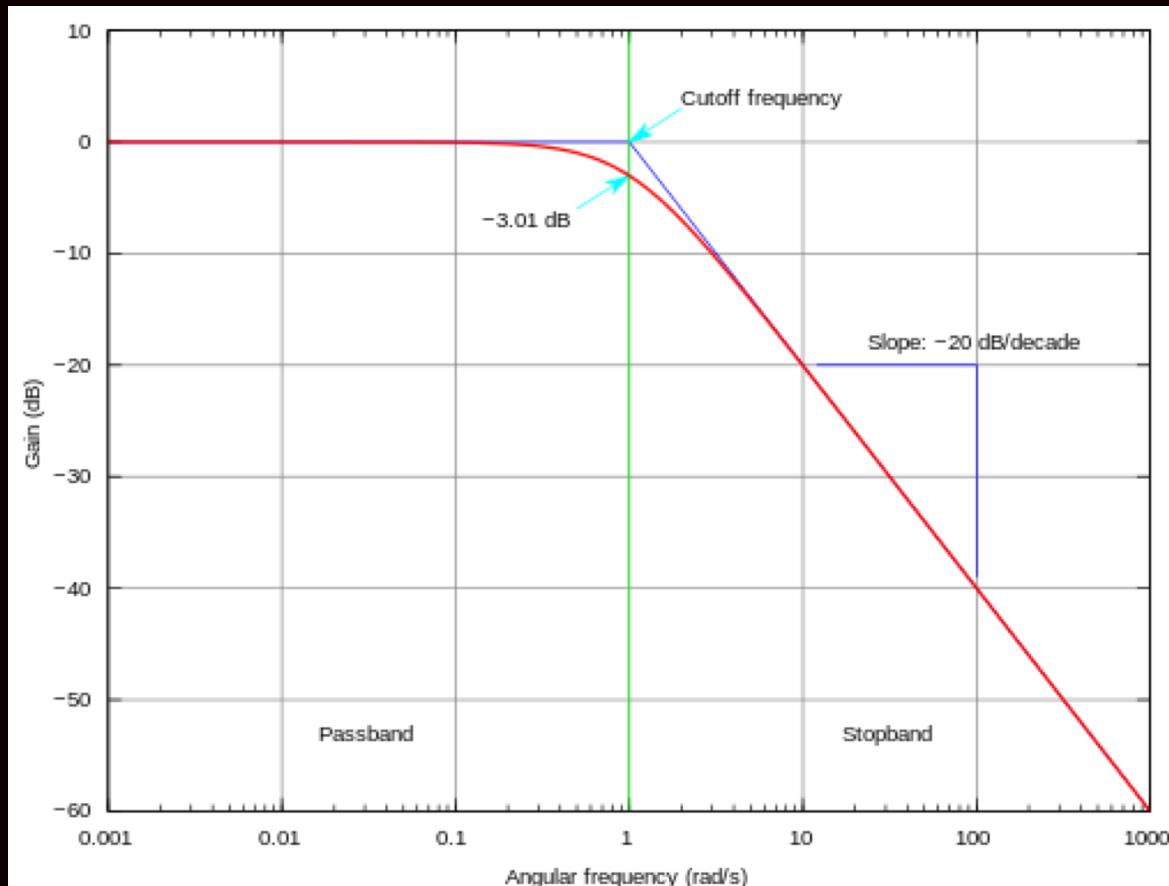
- Common cutoff slopes encountered are:

- 6 dB/oct
 - 12 dB/oct
 - 18 dB/oct
 - 24 dB/oct

For example: a cutoff slope of 6 dB/oct would imply that beyond the cutoff frequency, sound pressure will drop by 6 decibels for every octave increase in frequency

- Steeper slopes are not necessarily better:
 - Gentler slopes might be better for more subtle filtering
 - Steeper sloped filters can sound more dramatic and can draw attention to themselves

ANATOMY OF A CUTOFF SLOPE



- This is the response curve for a lowpass filter.
- Notice the filter has already started to remove frequencies below the cutoff point.
- The cutoff point is defined as the point at which 3dB of attenuation has occurred.

CSOUND FILTERS

- Csound possesses a vast range of filters

```
areson lowpass2 lowres lowresx lpf18 moogvcf moogladder reson  
resonr resonx resony resonz rezzy statevar svfilter tbvcf vlowres  
bqrez atone atonex tone tonex biquad biquada butterbp butterbr  
butterhp butterlp clfilt aresonk atonek lineto port portk resonk  
resonxk tlineto tonek dcblock dcblock2 pareq rbjeq eqfil nlfilt  
filter2 fofilter hilbert mode zfilter2
```

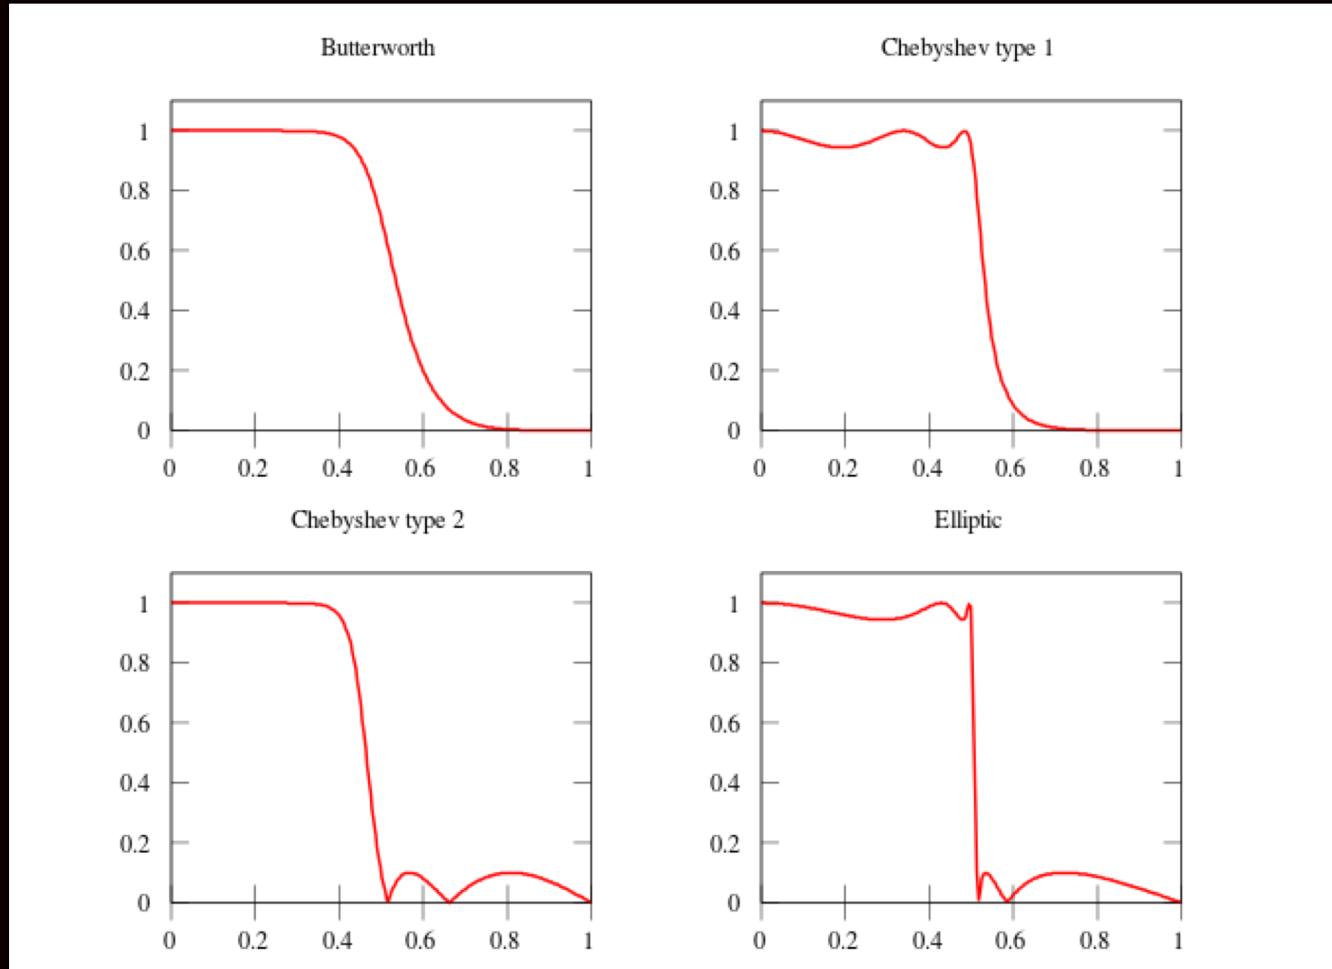
- A good place to start is the butterworth filters.

ares	butlp	asig, kcf	; lowpass filter
ares	buthp	asig, kcf	; highpass filter
ares	butbp	asig, kcf, kbw	; bandpass filter
ares	butbr	asig, kcf, kbw	; bandreject filter

- **kcf** = cutoff frequency
- **kbw** = bandwidth

SPECIFIC FILTER MODELS

- Different filters built upon different mathematical models (which often exist in both the analogue and digital domains) exhibit different characteristics. None are ideal.



By Alessio Damato - Own work, CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=427738>

- Butterworth offers a very flat passband so is often favoured in musical applications.

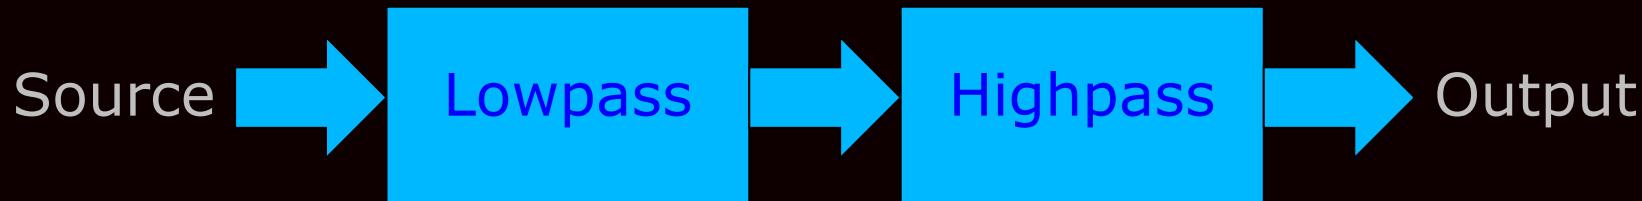
FILTERING EXAMPLE

- Lowpass filtering a sawtooth waveform

```
aSig  vco2  0.4, 220      ; create a sawtooth
kCF   expon 2000, p3, 50 ; create a filter cutoff envelope
aSig  butlp aSig, kCF    ; filter the sawtooth
      outs  aSig, aSig   ; send to the speakers
```

- **expon** is a variant of the **line** opcode in which the connecting segment is an exponential curve rather than a straight line. (Values should be alike in sign and the envelope should not cross zero.)
- It is acceptable to overwrite the input of an opcode with its output if we have no further use of the version provided by the input. This technique can simplify coding by leaving us with fewer variables to keep track of. It also saves RAM but this saving is fairly negligible on modern computers.
- **kCF** can also be a-rate; this might be a useful option if the cutoff will be changing at a high rate.

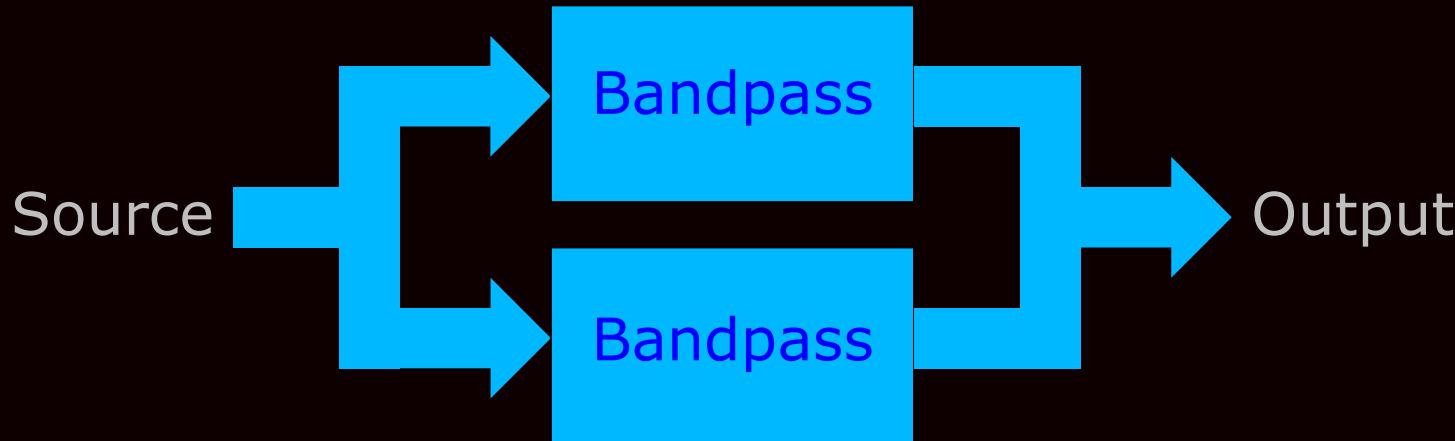
SERIES CONNECTION



```
aSig vco2 0.4, 200, 2, 0.5 ; square wave  
aSig butlp aSig, 5000 ; lowpass filter  
aSig buthp aSig, 2000 ; highpass filter  
outs aSig, aSig ; output
```

- Order of the filters doesn't matter.
- If the cutoff frequency of the highpass filter is above that of the lowpass filter then the signal can disappear completely.
- This arrangement can be used as a kind-of bandpass filter.

PARALLEL CONNECTION



```
aSig  vco2  0.4, 111, 2, 0.5 ; square wave
aBP1  butbp aSig, 1000, 333 ; bandpass filter 1
aBP2  butbp aSig, 2000, 667 ; bandpass filter 2
aMix =      aBP1 + aBP2      ; mix the two filters
          outs aMix, aMix       ; output
```

- Order of the filters doesn't matter.
- Note that we create two new audio variables as the outputs of the two bandpass filters.
- It is suggested to create another audio variable for the mix rather than mixing them at the output.

STATIC FILTERS vs FUNDAMENTAL-FOLLOWING FILTERS

- It is equally valid to design a subtractive synth within which a filter cutoff remains fixed or one which moves in ratio with changes in the fundamental frequency of the sound source.
- The reasons for choosing one or the other relate to the physics you might be emulating:
 - Filtering that is physically inate to the exciter (such as a guitar string) will shift in ratio with the source sound.
 - Filtering that is physically inate to a resonator (such as a guitar body) will remain fixed.
- See examples 5 and 6.



FUNDAMENTAL FOLLOWING FILTERS

- Also commonly referred to as 'key following'.

```
iCPS    cpsmidi                                ; read in MIDI notes  
aSig    vco2  0.4, iCPS                         ; create a sawtooth  
iBW     =      0.7 * iCPS                        ; bandwidth  
iRatio  =      5                                ; C.F ratio  
aSig    butbp aSig, iCPS*iratio, iBW ; filter the sawtooth  
       outs aSig, aSig                      ; send to the speakers
```

- In this example the bandpass filter will always accent the 5th harmonic partial, regardless of what note is played.
- It is also common practice to express the bandwidth of the bandpass filter in ratio with its frequency. This can be considered more 'musical' as the bandwidth control (in this example 0.7) will be expressing a fraction of an octave rather than a fixed number of cycles-per-second.
- It is also possible to devise a filter that is partially static and partially key-following by using a cutoff value that is an interpolation between a static and a key following value.